

Confidence Calibration under Ambiguous Ground Truth

Anonymous Authors

Abstract—Confidence calibration assumes a unique ground-truth label per input, an assumption that fails wherever annotators genuinely disagree. Post-hoc calibrators fitted on majority-voted labels can appear well-calibrated against voted labels yet remain substantially miscalibrated against the underlying annotator distribution. This failure is structural: under simplifying assumptions, Temperature Scaling is biased toward temperatures that underestimate annotator uncertainty, with true-label error increasing with annotation entropy. We propose ambiguity-aware post-hoc calibrators that optimize proper scoring rules against the full label distribution, with **Dirichlet-Soft** as the flagship method. For settings where only one-hot voted labels are available, **Label-Smooth Temperature Scaling (LS-TS)** constructs pseudo-soft targets from the model’s own confidence without annotator data. Across four benchmarks spanning real multi-annotator distributions (CIFAR-10H, ChaosNLI) and clinically-informed synthetic annotations (ISIC 2019, DermaMNIST), our methods substantially improve calibration performance evaluated by ECE, Brier score, and NLL; Dirichlet-Soft reduces true-label ECE by 59–86% relative to Temperature Scaling and achieves the best or tied-best Brier/NLL in 7 of 8 settings, while LS-TS reduces ECE by 7–78% without any annotator data.

Index Terms—Confidence calibration, ambiguous ground truth, annotator disagreement, temperature scaling, soft labels, proper scoring rules, label distribution learning.

1 INTRODUCTION

DEEP neural networks are increasingly deployed as decision-support tools in high-stakes domains. In medical imaging, classifiers assist dermatologists in diagnosing skin lesions and pathologists in identifying malignant tissue [1], [37]; in natural language processing, models flag toxic content, assess clinical notes, and support legal document review; in autonomous driving, perception systems make real-time safety decisions under uncertainty [2]. In each of these settings the model outputs not merely a predicted class but an implicit *reliability claim*: the confidence score attached to a prediction is understood by downstream systems and human operators as a direct probability that the prediction is correct. A radiologist who sees 90% confidence on a malignant-lesion classification may reduce the level of follow-up scrutiny; a content moderator acting on a 99% toxicity score may take irreversible action; an autonomous vehicle assigns braking decisions based on its detection confidence.

The faithfulness of stated confidence to empirical correctness — calibration — is therefore as important to deployment safety as accuracy itself [3], [4].

1.0.0.1 Confidence calibration.: This property is formalised as follows. A well-calibrated classifier $f : \mathcal{X} \rightarrow \Delta^K$ produces probability vectors $\hat{p}(x) = f(x)$ such that

$$\mathbb{P}(\hat{Y} = Y \mid \hat{p}_{\hat{Y}}(X) = p) = p, \quad \forall p \in [0, 1], \quad (1)$$

where $\hat{Y} = \arg \max_k \hat{p}_k(X)$ and Δ^K denotes the K -dimensional probability simplex. Guo et al. [6] showed that modern high-accuracy networks are systematically overconfident, and proposed Temperature Scaling (TS) — dividing logits by a scalar $T > 1$ before softmax — as an effective post-hoc remedy. Subsequent work extended post-hoc calibration to Platt scaling [8], isotonic regression [7], and Dirichlet calibration [9], and studied its behaviour across architectures [10] and tasks [11]. These methods share a common design principle: they fit a correction to the model’s logits using a held-out calibration set annotated with *one-hot labels*.

1.0.0.2 The hidden assumption.: The calibration condition (1) implicitly assumes that every input x has a unique ground-truth label Y . This assumption is violated whenever the task is inherently ambiguous or subjective. A skin lesion may be legitimately classified as malignant or benign by different dermatologists, depending on subtle visual features that experts genuinely disagree about [37]; a natural-language premise may or may not entail a hypothesis depending on pragmatic interpretation [15]; a low-resolution object may be categorically ambiguous even to careful human observers [14]. Such disagreement is not annotation *noise* to be corrected — it is irreducible *aleatoric* uncertainty inherent to the task. The correct calibration target is therefore not a single label Y but a *distribution over labels* $\pi(\cdot \mid x) \in \Delta^K$, which we call the *annotator distribution*, representing the probability that a randomly selected expert would assign each label to input x . Figure 1 illustrates three concrete instances of this phenomenon across image classification, natural language inference, and medical imaging.

1.0.0.3 The standard practice and its failure.: Standard practice aggregates annotations by majority vote to obtain a single *voted label* y^* and calibrates against it — and in many practical settings this one-hot label is the only supervision available. We show this is not merely a rough approximation but a *systematic failure*: calibrating

- Manuscript received XX XXX, XXXX; revised XX XXX, XXXX. This work was supported in part by [funding sources].
- The authors are with [institution]. E-mail: [email].




Fig. 1: **Left:** a CIFAR-10H image with dispersed human votes (cat/dog/bird), illustrating perceptual ambiguity in low-resolution vision. **Middle:** a ChaosNLI premise–hypothesis pair with split entailment/neutral judgments, illustrating semantic ambiguity. **Right:** an ISIC 2019 melanoma image paired with the clinician-informed label distribution used in our medical experiments, showing clinically plausible MEL/NV confusion. CIFAR-10H and ChaosNLI distributions are empirical human label distributions; the ISIC panel uses the dermatologist confusion model described in Appendix D.

to the voted label optimises the wrong objective, and this miscalibration worsens as genuine annotator disagreement increases. Strikingly, the most expressive voted-label calibrator (Dirichlet calibration with a full affine transform) consistently *underperforms* simple Temperature Scaling in true-label ECE — more flexibility with the wrong target makes things worse, not better. In a controlled experiment (Section 4), every standard voted-label calibrator fails to close the gap; the root cause is the calibration *target*, not the method’s capacity. This observation parallels the finding of Stutz et al. [5] for conformal prediction, where sets calibrated on voted labels systematically uncover the true annotator distribution.

1.0.0.4 Contributions.:

- 1) **Problem and theory** (Sections 3–5): we formalise true-label calibration under ambiguous ground truth and prove that Temperature Scaling is biased toward unduly low temperatures, with the miscalibration gap growing monotonically with annotation entropy.
- 2) **Methods** (Section 6): we propose a family of ambiguity-aware post-hoc calibrators — notably **Dirichlet-Soft** — that use the annotator distribution as the calibration target and require no model retraining; and **LS-TS**, an annotation-free variant that uses only voted labels yet substantially closes the calibration gap by constructing a data-driven soft pseudo-target.
- 3) **Evidence** (Sections 4 and 7): experiments across four benchmarks spanning vision (CIFAR-10H, ISIC 2019, DermaMNIST) and language (ChaosNLI) show that Dirichlet-Soft reduces ECE_{true} by 59–86% and LS-TS by 7–78% relative to TS, while standard calibrators can actively worsen true-label ECE.

2 RELATED WORK

2.0.0.1 Confidence calibration.: Guo et al. [6] established that modern high-accuracy networks are systematically overconfident and that Temperature Scaling (TS) is a

simple and effective post-hoc remedy. The multiclass calibration literature has since proposed isotonic regression [7], Platt scaling [8], Dirichlet calibration [9], and adaptive per-instance temperature variants [26]. Minderer et al. [10] showed that Vision Transformers are inherently better calibrated than CNNs, and Bai et al. [11] revisit when TS is and is not sufficient. Kumar et al. [?] propose verified uncertainty calibration through non-parametric testing. All of these methods share a common assumption: the calibration set contains one-hot labels, and calibration means aligning confidence with correctness relative to those labels. We show that this assumption — not the method’s architectural capacity — is the limiting factor when labels are genuinely distributional.

2.0.0.2 Soft labels, label smoothing, and knowledge distillation.: Training-time approaches have explored the benefits of soft targets. Label smoothing [12], [13] adds a fixed uniform component to one-hot targets and has been observed to improve both accuracy and calibration; however, it applies a dataset-global constant regardless of per-instance annotator disagreement. Knowledge distillation [25] uses teacher model outputs as soft training targets, improving student generalisation and calibration. Thulasidasan et al. [24] show that Mixup training substantially improves calibration by encouraging confident predictions only when training examples are well-separated. Collins et al. [?] collect CIFAR-10S, a companion to CIFAR-10H with individually elicited soft labels from every annotator, and demonstrate that training with per-annotator soft labels improves calibration over aggregated hard labels. Our work is the post-hoc calibration counterpart: we do not retrain the model, but instead re-target the calibration objective to the annotator distribution using only cached logits.

2.0.0.3 Annotation disagreement and multi-annotator learning.: Dense annotation datasets [14], [15] and multi-annotator learning methods [16], [17] have documented the ubiquity of label disagreement across domains. Aroyo and Welty [?] argue that disagreement is not noise but signal, and that discarding it by

majority aggregation loses information that is critical for understanding task difficulty. Gordon et al. [19] and Plank [20] highlight systematic epistemological problems with annotation aggregation in NLP, showing that models trained on majority labels inherit structural biases against minority annotator perspectives. Baan et al. [18] show that ECE is theoretically ill-defined when annotators genuinely disagree, because no unique “correct” label exists to serve as the reference; they propose instance-level uncertainty measures. We complement this critique by retaining ECE_{true} as a practical metric — it quantifies the overconfidence a user experiences when drawing a single label $\tilde{y} \sim \pi(\cdot|x)$ at deployment time, mirroring realistic usage — while supplementing it with Brier score and NLL, which are strictly proper scoring rules that directly penalise divergence from the full annotator distribution $\hat{\pi}$. Khurana et al. [21] study how annotator disagreement can inform selective abstention in NLP; unlike their approach, our framework targets all predictions and is applicable across vision and language without task-specific thresholding.

2.0.0.4 Proper scoring rules and distributional evaluation.: Gneiting and Raftery [39] provide the foundational theory of proper scoring rules: a loss is *strictly proper* if and only if it is uniquely minimized when the predicted distribution equals the true distribution. Brier score and NLL are canonical instances. This perspective motivates our choice of calibration objective: minimizing cross-entropy against $\hat{\pi}(x)$ is strictly proper over the predicted distribution $\hat{p}(x)$, while minimizing against a one-hot label is only proper with respect to a degenerate Dirac distribution and misidentifies the optimal calibrated output whenever $\hat{\pi}$ is non-degenerate. The proper scoring framework also clarifies why Dirichlet calibration with voted targets (Dirichlet-Hard) is not merely suboptimal but is consistent for the *wrong* estimand, explaining the empirical finding that it underperforms TS.

2.0.0.5 Conformal prediction under ambiguous ground truth.: Stutz et al. [5] demonstrate the set-prediction analogue of our finding: conformal prediction sets calibrated on voted labels systematically undercover the true annotator distribution, with coverage gaps of up to 10% on dermatology data. Their Monte Carlo conformal approach samples from the annotator distribution to compute coverage, closely paralleling our Monte Carlo ECE evaluation ($S = 100$ draws per example). Angelopoulos and Bates [?] provide foundational methodology for distribution-free uncertainty quantification, including the conformal framework that Stutz et al. build on. Our work addresses point-estimate calibration rather than set coverage: we produce well-calibrated scalar confidence scores, which conformal prediction does not. Calibration under covariate shift [22], [23] is orthogonal to our label-ambiguity setting and assumes a unique ground truth throughout.

2.0.0.6 Uncertainty quantification in high-stakes applications.: The deployment literature motivates the practical urgency of this problem. Kompa et al. [1] conduct a large-scale empirical study showing that medical image classifiers are consistently overconfident, even after calibration, and that this overconfidence degrades clinical decision quality. Ovadia et al. [3] demonstrate that popular uncertainty estimation methods (dropout, deep ensembles, variational inference) exhibit significantly degraded calibration under distribution

shift, highlighting that single-point calibration benchmarks may be overly optimistic. Feng et al. [2] survey uncertainty quantification in autonomous driving, where overconfident detections directly contribute to unsafe manoeuvres. These findings establish the applied context for our work: in each of these domains, label ambiguity is simultaneously a property of the task and an unaddressed source of calibration error. The present paper provides the first systematic post-hoc calibration framework that directly targets this source.

3 PROBLEM FORMULATION

3.1 Setup and Notation

Let $\mathcal{Y} = \{1, \dots, K\}$ be the label space. A classifier $f : \mathcal{X} \rightarrow \Delta^K$ outputs $\hat{p}(x) = \text{softmax}(z(x))$ from pre-softmax logits $z(x) \in \mathbb{R}^K$. The **annotator distribution** $\pi(\cdot | x) \in \Delta^K$ denotes the underlying ambiguous label distribution for input x and is estimated from m annotations as $\hat{\pi}_k(x) = \frac{1}{m} \sum_j \mathbf{1}[a_j = k]$. The **voted label** is $y^* = \arg \max_k \hat{\pi}_k(x)$. Standard calibration uses one-hot voted labels. For example, Temperature Scaling [6] minimizes

$$\mathcal{L}_{\text{TS}}(T) = -\frac{1}{n} \sum_{i=1}^n \log \text{softmax}(z_i/T)_{y_i^*}. \quad (2)$$

3.2 True-Label Calibration

Following Stutz et al. [5], we distinguish the voted label y^* from the **true label** $\tilde{y} \sim \pi(\cdot | x)$, drawn from the underlying ambiguous label distribution.

Definition 1 (True-Label Calibration). *f is **true-label calibrated** if for all $p \in [0, 1]$: $\mathbb{P}(\tilde{Y} = \hat{c}(X) | \hat{p}_{\hat{c}}(X) = p) = p$, where $\hat{c}(x) = \arg \max_k \hat{p}_k(x)$ and $\tilde{Y} \sim \pi(\cdot | X)$.*

The **voted-label ECE** $\text{ECE}_{\text{voted}}$ is the standard ECE against y^* . The **true-label ECE** ECE_{true} averages ECE over 100 draws $\tilde{y}_i \sim \pi(\cdot | x_i)$ ($B = 15$ equal-width bins).

4 MOTIVATING EXAMPLE

4.0.0.1 Setup.: We construct a 3-class 2D Gaussian dataset with three clusters: $x \sim \mathcal{N}((-3.2, 1.1), \text{diag}(0.60, 0.45))$ for class 0 with $\pi = [1, 0, 0]$; $x \sim \mathcal{N}((0, 0), \text{diag}(1.15, 0.75))$ for the ambiguous cluster with $\pi = [0, 0.70, 0.30]$; and $x \sim \mathcal{N}((3.2, -1.1), \text{diag}(0.60, 0.45))$ for class 2 with $\pi = [0, 0, 1]$. The voted label for the middle cluster is always class 1 because $0.70 > 0.30$. A 2-layer MLP (64 hidden units per layer) is trained on voted labels, and all calibrators are fitted on a held-out calibration split using the same voted labels.

4.0.0.2 Results.: Figure 2 shows that this mismatch is a structural consequence of the voted-label target. TS reduces $\text{ECE}_{\text{voted}}$ from 3.25% to 1.11%, but *increases* ECE_{true} from 6.26% to 9.64%. Platt scaling behaves similarly ($\text{ECE}_{\text{voted}} = 0.68\%$, $\text{ECE}_{\text{true}} = 9.70\%$), and Histogram Binning likewise ($\text{ECE}_{\text{voted}} = 1.32\%$, $\text{ECE}_{\text{true}} = 9.65\%$). Figure 2(f) further shows that this residual error is concentrated in the ambiguous cluster for all three methods. TS sets $T = 0.62 < 1$, boosting confidence toward the voted target and thereby pushing the model *further* from the true ambiguous label distribution. The failure is therefore not specific to TS:




Fig. 2: **Motivating example: all standard calibration methods fail.** (a) Toy dataset generation: only the middle Gaussian cluster is ambiguous, with $\pi = [0, 0.70, 0.30]$; orange points are drawn as label 1 (70%) and red points as label 2 (30%), while all receive the same voted label 1. (e) Summary of the toy results: all three voted-label calibrators (TS, Platt, Histogram Binning) lower ECE_{voted} but increase ECE_{true} , so voted-label evaluation masks the failure. (f) Stratified ECE_{true} for TS, Platt, and Histogram Binning: ambiguous examples are those from the middle Gaussian cluster (where annotators disagree); clear examples are those from the two unambiguous clusters (class 0 and class 2). The residual error is concentrated in ambiguous examples for all three methods.

standard post-hoc methods are useful for ECE_{voted} , but the voted-label target makes them ineffective for ECE_{true} .

5 THEORY: WHY STANDARD CALIBRATION FAILS

We now formalize, under simplifying assumptions, two complementary aspects of the failure of voted-label calibration under ambiguity: the *direction* of the TS bias (Proposition 1) and its *scaling* with label ambiguity (Proposition 2). Both propositions characterize *pointwise* miscalibration on individual examples; they describe when and why a voted-label calibrator misfits individual ambiguous inputs, rather than providing a theorem about population-level ECE. Figure 3 provides empirical support.

Proposition 1 (Direction of TS Bias). *Consider a calibration set containing an ambiguous cluster in which the voted label y^* equals the majority class for every example, but the true annotator probability of y^* is $\pi_{y^*}(x) = q < 1$. Under the assumption that the model is already reasonably accurate on this cluster ($\text{softmax}(z_i)_{y^*} > q$ for all i in the cluster), the voted-label loss exerts a downward pull on T for these examples, whereas a soft-label loss with target q exerts an upward pull. Consequently, $T_{TS} < T_{soft}^*$: TS selects a lower temperature than ambiguity-aware calibration requires. In a setting where the model is already overconfident ($\text{softmax}(z_i)_{y^*} \gg q$), this manifests as $T_{TS}^* < 1$; more generally, both optima may exceed 1, but the gap $T_{soft}^* - T_{TS}^* > 0$ persists, consistent with our experimental findings (e.g., $T_{TS} = 2.03$ vs. $T_{SLTS} = 3.18$ on CIFAR-10H ResNet-50).*

Proof sketch. On the ambiguous cluster, all voted labels are y^* , so the voted-label loss (2) is minimized by pushing $\text{softmax}(z_i/T)_{y^*} \rightarrow 1$, i.e. the loss decreases as T decreases. Hence the ambiguous examples exert a downward gradient on T . A soft-label loss with target $q < 1$ is instead minimized at $\text{softmax}(z_i/T^*)_{y^*} \approx q$, requiring a higher T^* whenever $\text{softmax}(z_i)_{y^*} > q$. The global TS optimum balances this downward pull against the upward pull from unambiguous

examples, but remains lower than the ambiguity-aware optimum. Full proof in Appendix .04. $\square$

Proposition 2 (True-Label Miscalibration and Annotation Entropy). *Let $H(x) = -\sum_k \pi_k(x) \log \pi_k(x)$ denote the annotation entropy of x . Assume (i) the model’s predicted confidence $\hat{p}_c(x)$ is approximately independent of $H(x)$ (similar accuracy across ambiguity levels), and (ii) $\pi_{y^*}(x)$ is a decreasing function of $H(x)$ (greater entropy implies lower majority-class probability). Under assumptions (i) and (ii), the expected per-example true-label calibration error of voted-label-calibrated predictions is non-decreasing in $H(x)$.*

Proof sketch. For unambiguous examples ($H(x) = 0$), $\pi(\cdot | x)$ is one-hot, so true-label and voted-label evaluation coincide and the per-example error is $|\hat{p}_c(x) - 1|$. As $H(x)$ increases, assumption (ii) gives $\pi_{y^*}(x) < 1$, introducing a gap between the voted-label target and the true annotator probability. Because TS is optimised against one-hot voted labels, it targets $\hat{p}_{y^*}(x) \approx 1$ for all examples regardless of their annotation entropy; hence $\hat{p}_c(x) > \pi_{y^*}(x)$ for ambiguous examples. The true-label calibration error is then $\hat{p}_c(x) - \pi_{y^*}(x)$, which by assumption (i) (fixed $\hat{p}_c$) increases as $\pi_{y^*}(x)$ falls with $H(x)$. $\square$

6 METHODS

All methods in this paper are *post-hoc*: they operate on cached logits from a fixed pre-trained model and require only a calibration set equipped with annotator distributions or individual annotations. Training-time approaches such as soft-label training [12], [13], mixup [24], and multi-annotator learning [16], [17] are orthogonal and not compared here.




Fig. 3: **Empirical validation of Proposition 2.** Test examples are grouped into equal-frequency bins by normalized annotation entropy $H(x)/\log K$. Each point shows the mean pointwise true-label calibration error $|\hat{p}_{\hat{e}}(x) - \pi_{\hat{e}}(x)|$ (error bars: ± 1 s.e.). TS error consistently increases with annotation entropy on CIFAR-10H, confirming Proposition 2. SLTS moderates but does not eliminate this trend.

6.1 Ambiguity-Aware Calibration Methods

6.1.0.1 Soft-Label Temperature Scaling (SLTS): SLTS minimizes the cross-entropy between the calibrated output and the empirical label distribution:

$$\begin{aligned} \mathcal{L}_{\text{SLTS}}(T) &= -\frac{1}{n} \sum_{i=1}^n \sum_{k=1}^K \hat{\pi}_k(x_i) \log \text{softmax}(z_i/T)_k \\ &= \frac{1}{n} \sum_{i=1}^n \text{KL}(\hat{\pi}(x_i) \parallel \text{softmax}(z_i/T)) + \text{const.} \quad (3) \end{aligned}$$

SLTS uses the same single-parameter family as TS but with a target that correctly reflects the annotator distribution.

Proposition 3 (Correctness of distributional target). *The per-example loss $\ell(q; \hat{\pi}) = -\sum_k \hat{\pi}_k \log q_k$ is a strictly proper scoring rule over $q \in \Delta^K$: it is uniquely minimised at $q = \hat{\pi}(x)$ for each x . Consequently, within the constrained family $\{\text{softmax}(z/T) : T > 0\}$, $\mathcal{L}_{\text{SLTS}}$ identifies the temperature T^* that minimises $\text{KL}(\hat{\pi} \parallel \text{softmax}(z/T^*))$, targeting the correct distributional objective, unlike $\mathcal{L}_{\text{TS}}$, which targets the one-hot voted label. Note that T^* need not achieve $\text{softmax}(z/T^*) = \hat{\pi}$ exactly within this one-parameter family; the proper scoring rule property guarantees correctness of the target, not achievability within the constrained family.*

6.1.0.2 Monte Carlo Temperature Scaling (MCTS): When individual annotation records are available rather than pre-aggregated label distributions, **MCTS** draws S samples $a_{is} \sim \hat{\pi}(x_i)$ per example and minimizes $\mathcal{L}_{\text{MCTS}}(T) = -\frac{1}{nS} \sum_{i,s} \log \text{softmax}(z_i/T)_{a_{is}}$. Formally, $\mathbb{E}_{\hat{y} \sim \hat{\pi}}[\text{CE}(\text{softmax}(z/T), \hat{y})] = \text{KL}(\hat{\pi} \parallel \text{softmax}(z/T)) + H(\hat{\pi})$, so MCTS converges to SLTS in expectation as $S \rightarrow \infty$. Table 1 shows that even $S = 1$ achieves ECE_{true} of 1.53%, matching SLTS within rounding error: any dataset annotated even once by a randomly selected annotator is immediately amenable to strong ambiguity-aware calibration (see Appendix .0.10 for full analysis).

6.1.0.3 Vector Scaling (VS): Extends SLTS with per-class temperatures $\mathbf{T} = (T_1, \dots, T_K)$, accommodating

TABLE 1: **MCTS convergence on CIFAR-10H (ResNet-50).** Mean $\pm$ std over 5 seeds; S = MC annotation samples per calibration example.

Method	ECE (%)	T^*
TS (voted labels)	4.29	2.030
MCTS $S = 1$	1.53 ± 0.01	3.174 ± 0.041
MCTS $S = 5$	1.53 ± 0.02	3.166 ± 0.014
MCTS $S = 20$	1.52 ± 0.00	3.187 ± 0.013
MCTS $S = 50$	1.52 ± 0.01	3.185 ± 0.007
MCTS $S = 200$	1.52 ± 0.00	3.178 ± 0.003
SLTS ($S \rightarrow \infty$)	1.51	3.180

datasets where some classes are systematically more ambiguous than others.

6.1.0.4 Distributional Isotonic Regression (IR-Soft): Fits a monotone mapping from predicted confidence to mean annotator accuracy via the Pool Adjacent Violators Algorithm (PAVA), using the empirical label distribution as target.

6.1.0.5 SoftPlatt: The diagonal affine transform $W \cdot z + b$ (same structure as Platt) fitted against the annotator distribution via KL divergence. This isolates the contribution of the distributional target from any architectural difference.

6.1.0.6 Dirichlet-Soft: The full $K \times K$ affine $Wz + b$ fitted by minimizing $\text{KL}(\hat{\pi}(x) \parallel \text{softmax}(Wz + b))$ with ODIR regularization ($\lambda = 10^{-3}$) [9]. This is the most expressive parametric ambiguity-aware calibrator.

6.2 Annotation-Free Calibration

The methods above require annotator distributions $\hat{\pi}(x_i)$ at calibration time. When only voted labels y_i^* are available, we propose **Label-Smooth Temperature Scaling (LS-TS)**, which constructs a pseudo-target distribution from the model’s own predictions without any additional annotations.

Let $\bar{\varepsilon} = \frac{1}{n} \sum_{i=1}^n (1 - \hat{p}_{y_i^*}(x_i))$ be the mean complement of the model’s voted-class confidence on the calibration set. Define the pseudo-target distribution

$$\tilde{\pi}_i^{\text{LS}} = (1 - \bar{\varepsilon}) e_{y_i^*} + \frac{\bar{\varepsilon}}{K} \mathbf{1}, \quad (4)$$

and minimize $\mathcal{L}_{\text{SLTS}}$ (Eq. 3) with these targets. The global smoothing weight $\bar{\varepsilon}$ is a data-driven estimate of average annotator disagreement: if the model is 80% confident on the voted class, the remaining 20% is spread uniformly across all K classes. The result is $T_{\text{LS}}^* > T_{\text{TS}}^*$ whenever $\bar{\varepsilon} > 0$, moving the calibrator in the correct direction even without annotator data.

6.2.0.1 Theoretical interpretation.: The smoothing weight $\bar{\varepsilon}$ satisfies a moment-matching fixed point: the expected pseudo-label mass on the voted class equals the average model confidence, $\mathbb{E}[\pi_{y^*}^{\text{LS}}] = \mathbb{E}[\hat{p}_{y^*}]$. Algorithmically, LS-TS can be viewed as a single E-step of an EM algorithm in which the latent annotator distribution is estimated from the current (pre-calibration) model, followed by one M-step that optimises T given those pseudo-labels. This connection to knowledge distillation [25] explains the self-referential nature of the approach: the uncalibrated model both supplies the pseudo-targets and is then calibrated against them. Section 7.7 compares LS-TS against three simpler smoothing strategies and confirms that the data-driven global $\bar{\varepsilon}$ is the key design choice. Appendix G further shows that even an *adaptive* voted-label calibrator (one that learns a per-instance temperature from logit-derived features) fails to improve over TS, reinforcing that the target distribution, not the flexibility of the calibration map, is the limiting factor.

7 EXPERIMENTS

We evaluate on four benchmarks with multi-annotator data, each with two backbones: CIFAR-10H (ResNet-50, ViT-B/16), ChaosNLI (RoBERTa-Large, DeBERTa-v3), ISIC 2019 (EfficientNet-B4, ViT-S/16), and DermaMNIST (ResNet-18, ViT-S/16).

7.0.0.1 Baselines.: We compare against (i) **Uncalibrated**: raw softmax; (ii) **TS**: Temperature Scaling on voted labels [6]; (iii) **ATS** (Adaptive Temperature Scaling [26]): per-instance temperature predicted from four logit-derived features ($\max_k z_k$, prediction entropy, top-2 margin, top-1 confidence) via a linear model, trained with the same voted-label NLL as TS (the most direct adaptive extension of TS without annotator data); (iv) **Platt (PS)**: per-class weight and bias on logits, calibrated against voted labels [8], [9]; (v) **Dirichlet-Hard**: Dirichlet calibration [9] with the full $K \times K$ weight matrix and bias, calibrated against voted labels. Dirichlet-Hard uses the *same architecture* as our Dirichlet-Soft but with voted-label targets; the contrast isolates the effect of the calibration target from model capacity. ATS similarly tests whether richer architecture, without changing the target, is sufficient to fix the calibration gap.

7.0.0.2 Metrics.: We report two families of metrics, which differ in how the ground-truth label is treated:

Binning-based ECE metrics (ECE_{true} , aECE, cwECE) are estimated by drawing $S = 100$ labels $\hat{y}_i^{(s)} \sim \hat{\pi}(x_i)$ per example and averaging calibration error over draws. ECE_{true} uses $B = 15$ equal-width confidence bins; aECE (adaptive ECE [27]) uses equal-mass bins that adapt to the confidence distribution; cwECE [9] averages per-class ECE over K classes. These mirror the evaluation protocol faced by a practitioner who observes one label per prediction at test time.

Soft-target metrics (Brier score Br, NLL) are computed directly against the soft target $\hat{\pi}$:

$$\begin{aligned} \text{Br}_i &= \|\hat{p}(x_i) - \hat{\pi}(x_i)\|^2, \\ \text{NLL}_i &= -\sum_k \hat{\pi}_k(x_i) \log \hat{p}_k(x_i) = H(\hat{\pi}(x_i), \hat{p}(x_i)), \end{aligned} \quad (5)$$

averaged over the test set. Brier and NLL are strictly proper scoring rules that directly penalise divergence from the full annotator distribution, complementing the binning-based ECE metrics.

We abbreviate ECE_{true} , aECE, cwECE, Br, NLL throughout.

7.1 CIFAR-10H: Image Classification with Human Label Noise

7.1.0.1 Dataset and models.: CIFAR-10H [14] provides ≈ 51 human annotations per image for the 10 000-image CIFAR-10 test set, directly exposing perceptual ambiguity in low-resolution natural images through repeated human labeling. The test set is split into calibration ($n = 5,000$) and evaluation (5,000) by stratified sampling (seed 42). We evaluate two architectures: **ResNet-50** [28] pretrained on ImageNet-1k, fine-tuned for 30 epochs (AdamW, cosine annealing), achieving 97.3% top-1 accuracy. **ViT-B/16** [29] pretrained on ImageNet-21k, fine-tuned with the same protocol, achieving 98.1% top-1 accuracy.

7.1.0.2 Results.: Both CIFAR-10H backbones show the same pattern: hard-label baselines still leave 4–5% ECE_{true} , while ambiguity-aware methods reduce it to below 2%. Notably, ATS (adaptive per-instance temperature) achieves 4.10% ECE on both architectures, no better than global TS (4.29% / 4.48%) and far behind LS-TS (1.57% / 2.37%), confirming that per-instance temperature adaptation cannot compensate for the wrong calibration target. Dirichlet-Soft achieves the best Brier score and NLL on both architectures (0.109/0.271 on ResNet-50; 0.100/0.255 on ViT-B/16), confirming that expressive distribution-aware calibrators provide the best distributional fit. LS-TS, using only voted labels, already reduces ECE from 4.29% to 1.57% (ResNet-50) and from 4.48% to 2.37% (ViT-B/16), competitive with fully ambiguity-aware methods.

7.2 ChaosNLI: Natural Language Inference

We next evaluate on **ChaosNLI** [15]: 100 human annotations per example for the combined SNLI+MNLI subset ($n = 3,113$), where disagreement reflects genuine semantic ambiguity rather than annotation noise. We split the data into calibration and test sets (50/50, stratified by majority-vote label to preserve class balance). We evaluate **RoBERTa-Large** [30] and **DeBERTa-v3-base** [31], both fine-tuned on MNLI training data; ChaosNLI draws from the SNLI and MNLI dev/test splits, so no training examples are re-used at evaluation.

7.2.0.1 Results.: The same pattern holds in NLI: hard-label baselines leave 10–12% ECE, and ATS (11.36% on RoBERTa-L, 13.09% on DeBERTa-v3) performs *worse* than global TS on DeBERTa-v3, further demonstrating that adaptive capacity without the correct target is ineffective. The best distributional fit is achieved by Dirichlet-Soft and

TABLE 2: **CIFAR-10H and ChaosNLI results.** ECE/aECE/cwECE averaged over 100 sampled labels $\tilde{y} \sim \hat{\pi}$; Brier/NLL computed against soft target $\hat{\pi}$ (Eq. 5). Best per architecture in **bold**.

Method	ResNet-50						ViT-B/16					
	T	ECE↓	aECE↓	cwECE↓	Br↓	NLL↓	T	ECE↓	aECE↓	cwECE↓	Br↓	NLL↓
CIFAR-10H (≈ 51 annotations per image, $K=10$)												
<i>Voted-label baselines</i>												
Uncalibrated	—	4.97	5.71	1.09	0.120	0.692	—	4.99	5.36	1.06	0.111	0.678
TS	2.03	4.29	4.25	0.91	0.112	0.363	2.04	4.48	4.40	0.93	0.106	0.346
ATS	2.33	4.10	4.01	0.87	0.113	0.334	2.26	4.10	4.04	0.86	0.105	0.319
Platt	—	4.29	4.23	0.91	0.112	0.372	—	4.54	4.38	0.93	0.105	0.363
Dirichlet-Hard	—	4.46	4.44	0.95	0.114	0.395	—	4.70	4.62	0.97	0.107	0.394
<i>Ambiguity-aware (ours)</i>												
SLTS	3.18	1.51	1.39	0.45	0.110	0.293	3.07	0.85	0.56	0.39	0.102	0.278
SoftPlatt	—	1.52	1.31	0.35	0.110	0.288	—	0.88	0.47	0.24	0.101	0.272
VS	—	1.35	1.26	0.39	0.109	0.289	—	0.92	0.50	0.32	0.101	0.274
IR-Soft	—	0.72	0.83	0.45	0.115	0.340	—	0.91	0.61	0.43	0.105	0.321
Dirichlet-Soft	—	1.25	1.06	0.35	0.109	0.271	—	0.72	0.41	0.25	0.100	0.255
<i>Annotation-free (ours)</i>												
LS-TS	3.08	1.57	1.31	0.46	0.109	0.293	2.76	2.37	2.21	0.53	0.103	0.285
Method	RoBERTa-Large						DeBERTa-v3					
	T	ECE↓	aECE↓	cwECE↓	Br↓	NLL↓	T	ECE↓	aECE↓	cwECE↓	Br↓	NLL↓
ChaosNLI (SNLI+MNLI, 100 annotations per example, $K=3$)												
<i>Voted-label baselines</i>												
Uncalibrated	—	27.79	27.77	18.61	0.650	1.384	—	35.97	35.92	24.02	0.743	2.148
TS	2.42	10.55	10.38	7.30	0.537	0.886	3.81	11.63	11.57	8.48	0.549	0.900
ATS	2.28	11.36	11.29	7.73	0.538	0.883	3.57	13.09	13.07	9.06	0.552	0.900
Platt	—	11.16	10.92	7.28	0.530	0.875	—	12.43	12.36	8.43	0.539	0.896
Dirichlet-Hard	—	11.55	11.40	7.64	0.535	0.889	—	12.69	12.63	8.84	0.539	0.895
<i>Ambiguity-aware (ours)</i>												
SLTS	3.41	3.22	2.59	4.62	0.520	0.862	5.28	3.45	3.36	6.53	0.529	0.877
SoftPlatt	—	2.66	2.32	2.43	0.509	0.839	—	3.17	2.92	3.18	0.511	0.841
VS	—	3.46	3.07	5.13	0.518	0.857	—	3.91	3.35	6.46	0.525	0.869
IR-Soft	—	2.65	2.59	6.24	0.549	0.934	—	2.15	3.96	7.20	0.554	0.942
Dirichlet-Soft	—	2.57	1.71	2.03	0.510	0.839	—	3.19	2.71	2.95	0.509	0.836
<i>Annotation-free (ours)</i>												
LS-TS	4.40	4.16	3.72	5.61	0.522	0.873	6.15	2.65	2.99	6.32	0.529	0.881

SoftPlatt: on RoBERTa-L, SoftPlatt achieves the best Brier (0.509 vs. Dir-Soft 0.510) and both tie on NLL (0.839); on DeBERTa-v3, Dirichlet-Soft achieves the best Brier (0.509) and NLL (0.836). LS-TS reduces ECE from 10.55% to 4.16% (RoBERTa-L) and from 11.63% to 2.65% (DeBERTa-v3) using only voted labels, outperforming all hard-label baselines by a wide margin.

7.2.0.2 ATS vs. LS-TS: target matters more than architecture.: Table 3 confirms that ATS fails to improve over TS in all 8 settings and in several cases degrades it (ECE rises from 11.63% to 13.09% on DeBERTa-v3; from 16.59% to 17.53% on ISIC ViT-S/16). The per-instance temperature learned from logit features adapts to model uncertainty, not to annotator disagreement. In contrast, LS-TS, which uses the same global-temperature architecture as TS but shifts the calibration target toward a smoothed distribution, reduces ECE by 7–78% across all 8 settings without any annotator

data (see Appendix G for full analysis).

7.3 ISIC 2019: Skin Disease Diagnosis

We evaluate on **ISIC 2019** [32], [33] (8 skin conditions, $\approx 25,000$ images), a clinically realistic differential-diagnosis task in which several lesion categories have overlapping visual morphology. We use $m = 9$ synthetic dermatologist annotations per image, generated by sampling from a clinically calibrated 8×8 confusion matrix following Dawid and Skene [34] (Appendix D), with diagonal agreement matched to Liu et al. [35]: mean 75%. Because no per-image reader study is publicly available for this dataset, annotation ambiguity is modelled *class-conditionally*: each image’s confusion depends on its consensus class, not on its individual visual content, which is a deliberate simplification of true instance-level annotator variation. We evaluate **EfficientNet-B4** [36] and **ViT-S/16** [29] on a stratified 70/15/15 split.




Fig. 4: Reliability diagrams for four representative methods — CIFAR-10H ResNet-50 (ECE_{true}). Red shading indicates overconfidence; green indicates underconfidence. Uncal and TS remain overconfident; SLTS substantially corrects this with a single temperature parameter; Dirichlet-Soft reduces the residual gap further. Reliability diagrams for all remaining methods are provided in Appendix .014.




Fig. 5: Reliability diagrams for four representative methods — ChaosNLI RoBERTa-Large (ECE_{true}). Same method selection as Figure 4. Red shading indicates overconfidence; green indicates underconfidence. NLI models are severely overconfident before calibration; ambiguity-aware methods substantially correct this.

TABLE 3: ATS vs. TS and LS-TS across all 8 settings (annotation-free methods only). ECE_{true} (%) and Brier. ATS predicts a per-instance temperature from logit features but is trained with the same voted-label NLL as TS. **Bold** = best annotation-free method per column.

	C10H R50		C10H ViT		NLI RoBERTa		NLI DeBERTa		ISIC ENet		ISIC ViT		Derm R18		Derm ViT	
Method	ECE	Br	ECE	Br	ECE	Br	ECE	Br	ECE	Br	ECE	Br	ECE	Br	ECE	Br
TS	3.90	.113	4.01	.104	10.55	.537	11.63	.549	18.54	.560	16.59	.617	23.05	.686	24.36	.683
ATS	4.10	.113	4.10	.105	11.36	.538	13.09	.552	18.83	.560	17.53	.631	22.82	.687	25.35	.686
LS-TS	1.49	.111	1.62	.102	4.16	.522	2.65	.529	9.34	.528	15.49	.619	5.05	.630	7.36	.615

7.3.0.1 ISIC 2019 results.: The same pattern holds under synthetic annotators: Dirichlet-Soft achieves the best Brier and NLL on both backbones (0.518/1.084 on ENet-B4; 0.563/1.175 on ViT-S/16), while IR-Soft achieves the lowest ECE. LS-TS halves the ECE of TS on ENet-B4 (18.54%→9.34%) using only voted labels; the improvement is smaller on ViT-S/16 (16.59%→15.49%), where model confidence is a weaker proxy for annotation ambiguity.

7.3.0.2 Annotator confusion model.: Figure 33 shows the clinically calibrated 8×8 confusion matrix used to generate synthetic annotations for ISIC 2019. The MEL/NV pair (melanoma/nevi, orange dashed) is the most consequential and most confused: diagonal rates of 73%/76% with

14%/15% cross-confusion, reflecting the well-documented difficulty of distinguishing early melanoma from benign nevi [37]. The AK/BKL/SCC cluster (purple dotted) forms a second high-confusion group (diagonal 65–67%), capturing clinically significant overlap between actinic keratosis and squamous cell carcinoma. The mean diagonal agreement $\bar{C}_{ii} = 75\%$ matches the inter-reader agreement reported by Liu et al. [35]. Full matrix values and construction details are given in Appendix D.

TABLE 4: **ISIC 2019 and DermaMNIST results.** ECE/aECE/cwECE averaged over 100 sampled labels $\tilde{y} \sim \hat{\pi}$; Brier/NLL computed against soft target $\hat{\pi}$ (Eq. 5). Best per architecture in **bold**.

Method	EfficientNet-B4						ViT-S/16					
	T	ECE↓	aECE↓	cwECE↓	Br↓	NLL↓	T	ECE↓	aECE↓	cwECE↓	Br↓	NLL↓
ISIC 2019 ($m=9$ synthetic annotators, mean diagonal agreement 75%)												
<i>Voted-label baselines</i>												
Uncalibrated	—	25.76	25.84	6.67	0.600	2.710	—	23.31	23.33	6.55	0.651	1.978
TS	1.79	18.54	18.49	5.02	0.560	1.653	1.42	16.59	16.60	5.23	0.617	1.563
ATS	1.71	18.83	18.80	5.03	0.560	1.733	1.12	17.53	17.48	5.46	0.631	1.730
Platt	—	19.89	19.81	5.24	0.562	1.659	—	17.19	17.17	4.74	0.597	1.583
Dirichlet-Hard	—	20.36	20.32	5.33	0.563	1.674	—	17.76	17.72	4.86	0.600	1.597
<i>Ambiguity-aware (ours)</i>												
SLTS	4.46	9.71	9.74	2.92	0.528	1.146	3.15	6.96	6.68	3.29	0.586	1.233
SoftPlatt	—	9.25	9.31	2.23	0.523	1.108	—	8.05	7.90	2.05	0.565	1.190
VS	—	9.03	9.11	2.42	0.524	1.121	—	8.50	8.59	2.24	0.567	1.198
IR-Soft	—	1.75	2.43	4.28	0.538	1.284	—	1.73	1.59	5.78	0.621	1.466
Dirichlet-Soft	—	7.50	7.48	2.00	0.518	1.084	—	6.85	6.80	1.77	0.563	1.175
<i>Annotation-free (ours)</i>												
LS-TS	4.30	9.34	9.46	2.98	0.528	1.149	3.98	15.49	15.47	4.52	0.619	1.303
Method	ResNet-18						ViT-S/16					
	T	ECE↓	aECE↓	cwECE↓	Br↓	NLL↓	T	ECE↓	aECE↓	cwECE↓	Br↓	NLL↓
DermaMNIST ($m=5$ synthetic annotators, overall agreement 64.7%)												
<i>Voted-label baselines</i>												
Uncalibrated	—	34.35	34.26	10.15	0.767	2.689	—	37.12	37.10	10.82	0.786	3.766
TS	1.86	23.05	22.82	6.97	0.686	1.654	2.60	24.36	24.26	7.29	0.683	1.664
ATS	1.86	22.82	22.75	6.97	0.687	1.606	2.53	25.35	25.24	7.53	0.686	1.650
Platt	—	23.64	23.56	7.08	0.682	1.667	—	23.82	23.59	7.10	0.675	1.681
Dirichlet-Hard	—	24.58	24.46	7.33	0.690	1.771	—	24.84	24.68	7.49	0.682	1.897
<i>Ambiguity-aware (ours)</i>												
SLTS	3.51	4.61	4.44	3.56	0.622	1.376	5.11	3.58	3.06	3.37	0.610	1.350
SoftPlatt	—	4.19	3.55	1.61	0.612	1.312	—	3.56	2.93	1.29	0.601	1.289
VS	—	4.45	3.83	3.23	0.621	1.363	—	4.03	3.71	2.84	0.611	1.331
IR-Soft	—	2.15	2.34	4.55	0.634	1.440	—	3.79	3.73	4.59	0.628	1.431
Dirichlet-Soft	—	3.55	3.36	1.22	0.610	1.305	—	3.32	2.87	1.14	0.601	1.286
<i>Annotation-free (ours)</i>												
LS-TS	3.03	5.05	4.82	3.79	0.630	1.394	4.19	7.36	7.06	3.68	0.615	1.362

7.4 DermaMNIST: Supplementary Skin Disease Experiment

We additionally evaluate on **DermaMNIST** [38] (7-class, derived from HAM10000 [32]; $m = 5$ synthetic annotators; overall agreement $\approx 64.7\%$) with ResNet-18 and ViT-S/16. TS leaves ECE = 23.1% on ResNet-18; ambiguity-aware methods reduce it to 4.6% (SLTS), 2.2% (IR-Soft), and 3.6% (Dirichlet-Soft), with Dirichlet-Soft also achieving the best cwECE (1.2%), Brier (0.610), and NLL (1.305). ViT-S/16 follows the same pattern (Dirichlet-Soft: 3.3%/0.601/1.286). Both backbones show $T_{\text{SLTS}} > T_{\text{TS}} > 1$, consistent with Proposition 1.

DermaMNIST is the most ambiguous benchmark in our study, with mean annotator agreement of 64.7% compared to 75% for ISIC 2019, $\approx 80\%$ for CIFAR-10H, and $\approx 67\%$ for ChaosNLI. Correspondingly, TS ECE_{true} (23.1%) is the highest across all benchmarks, reflecting the widest gap

between the voted-label target and the true annotator distribution. The Dirichlet-Hard calibrator, despite being the most expressive voted-label method (full $K \times K$ affine), worsens ECE relative to TS on both backbones (24.6% vs. 23.1% on ResNet-18; 24.8% vs. 24.4% on ViT-S/16), again confirming that expressive architecture cannot compensate for a miscalibrated target. LS-TS reduces ECE by 78% relative to TS on ResNet-18 (23.1% $\rightarrow$ 5.1%) — the largest relative improvement of the annotation-free method across all settings — because the high level of ambiguity means model confidence (complement 35.3%) correlates well with the fraction of annotators choosing the minority class.

7.5 Cross-Benchmark Analysis

Three consistent patterns emerge across all four benchmarks. **(1) Target, not capacity, is the bottleneck.** Dirichlet-Hard — the most expressive voted-label calibrator — consistently




Fig. 6: **Inter-reader confusion matrix for ISIC 2019.** Entry C_{ij} is the probability that a dermatologist labels an image as class j given consensus label i . Orange dashed box: MEL/NV most-confused pair (73–76% diagonal). Purple dotted: AK/BKL/SCC high-confusion cluster. Mean diagonal agreement $\bar{C}_{ii} = 75\%$.

underperforms simple TS in both ECE and Brier (e.g., 4.46% vs. 4.29% on CIFAR-10H R50; 11.55% vs. 10.55% on ChaosNLI RoBERTa-L; see Tables 2 and 4). This rules out architectural capacity as the bottleneck and directly supports the claim that calibration target is limiting. **(2) Ambiguity-aware methods consistently dominate.** Soft-target methods improve ECE by 59–86% relative to TS on the two natural-annotation benchmarks (CIFAR-10H and ChaosNLI) and by 72–81% on the medical imaging benchmarks. Dirichlet-Soft achieves the best Brier score and NLL in 7 of 8 settings, making it the recommended default whenever annotator distributions are available. **(3) LS-TS is effective but dataset-dependent.** LS-TS substantially reduces ECE without any annotator data in all settings except ISIC ViT-S/16 (where $T_{LS} < T_{SLTS}$, indicating model confidence under-estimates annotation ambiguity for this architecture–dataset pair). This divergence highlights the assumption underlying LS-TS: that model uncertainty is a reliable proxy for annotation disagreement — an assumption that holds broadly but can fail when the pre-softmax logit magnitude is not well correlated with the fraction of annotators that choose the minority class.

7.6 Reliability Diagrams

Figures 4–5 show reliability diagrams for CIFAR-10H and ChaosNLI; Figures 7–8 extend the comparison to the medical imaging benchmarks. Across all four datasets, Uncal and TS are systematically overconfident; SLTS substantially corrects this with a single temperature parameter; Dirichlet-Soft closes the residual gap. Full 3×5 panels covering all 13 methods and 8 architectures are in Appendix .014.

7.7 Ablation: LS-TS Smoothing Strategies

Table 5 compares LS-TS against three simpler annotation-free smoothing strategies on CIFAR-10H and ChaosNLI. All

methods use the same single-temperature family; only the pseudo-target construction differs.

Fixed-LS with $\varepsilon=0.1$ consistently *underperforms* or even *degrades* below TS (e.g., ECE 7.44% vs. 3.90% on CIFAR-10H ResNet-50), confirming that an arbitrary fixed smoothing weight is insufficient. Ent-LS (per-instance entropy) is inconsistent across settings: it collapses on DeBERTa-v3 (8.58%), suggesting prediction entropy is a noisy proxy when models are strongly confident. CC-LS (per-class ε) nearly matches LS-TS on all four settings (within 0.1 pp ECE), confirming that the global mean is a good summary of per-class behaviour; the slight advantage of CC-LS on DeBERTa-v3 is within single-seed variance.

7.8 Robustness

Table 6 and Figure 9 summarise three robustness dimensions. **Multi-seed stability:** standard deviations of ECE_{true} over 5 random seeds are uniformly $2 \times$ – $10 \times$ smaller than the ECE gaps between methods, confirming that rankings are not artifacts of a particular data split or annotation draw; full bar charts are in Appendix .014. **Calibration set size:** ambiguity-aware methods stabilise below 2% ECE on CIFAR-10H using as few as 5–10% of the calibration data (≈ 250 examples); full results are in Appendix .012. **Annotation count:** ambiguity-aware calibration is robust once $m \geq 2$ annotations per image — going from 2 to 51 annotations reduces SLTS ECE_{true} by less than 0.3 pp — while TS ECE_{true} *increases* with m , because the annotator distribution diverges further from the one-hot voted label as consensus strengthens; full annotation-count results are in Appendix .013.

8 DISCUSSION

8.0.0.1 Sacrificing voted-label calibration.: A natural question is whether ambiguity-aware methods degrade ECE_{voted} . This is expected behaviour: voted-label calibration is epistemically unreliable when annotators genuinely disagree, as the voted label overstates the majority-class probability. Practitioners who must satisfy a hard ECE_{voted} constraint can apply TS on top of Dirichlet-Soft or SLTS outputs.

8.0.0.2 Practical guidelines.: Use **Dirichlet-Soft** as the default when annotator distributions are available: it achieves the best or near-best Brier score and NLL on most benchmarks (best NLL in 7/8 settings; best or tied Brier in 7/8 settings) while maintaining competitive ECE; SoftPlatt is competitive on Brier for NLI tasks. Use **SLTS** when a single-parameter method is preferred for simplicity; **IR-Soft** when the lowest ECE is the sole criterion; and **LS-TS** when the calibration set contains only one-hot voted labels.

8.0.0.3 Connection to proper scoring rules and distributional calibration.: The shift from one-hot to soft targets is not merely a methodological convenience: it corresponds to replacing an *improper* calibration criterion with a *proper* one. As Proposition 3 makes explicit, minimizing the cross-entropy against $\hat{\pi}$ is equivalent to minimizing $KL(\hat{\pi} \parallel \hat{p})$, a divergence that is uniquely minimized at $\hat{p} = \hat{\pi}$. This places our framework within the broader theory of proper scoring rules [39], where the Brier score and NLL play the role of strictly proper objectives. Voted-label calibration




Fig. 7: **Reliability diagrams — ISIC 2019 EfficientNet-B4 (ECE_{true})**. Red shading = overconfidence; green = underconfidence. Hard-label baselines remain severely overconfident at 18–26% ECE; ambiguity-aware methods substantially correct this, with IR-Soft achieving the lowest ECE.




Fig. 8: **Reliability diagrams — DermaMNIST ResNet-18 (ECE_{true})**. Same layout as Figure 4. TS leaves 23% ECE; Dirichlet-Soft reduces it to 3.6% while achieving the best Brier/NLL.

TABLE 5: **LS-TS ablation: annotation-free smoothing strategies**. ECE_{true} (%), Brier score, and NLL on CIFAR-10H and ChaosNLI. All methods use the same single-temperature family; only the pseudo-target construction differs. Fixed-LS uses $\varepsilon=0.1$ regardless of the data; Ent-LS uses per-instance entropy; CC-LS uses per-class mean complement-confidence. Best annotation-free method per column in **bold**.

Method	CIFAR R50			CIFAR ViT			Chaos RoBERTa			Chaos DeBERTa		
	ECE	Br	NLL	ECE	Br	NLL	ECE	Br	NLL	ECE	Br	NLL
TS (voted)	3.90	.1126	.346	4.01	.1043	.320	10.55	.537	.886	11.63	.549	.900
Fixed-LS ($\varepsilon=0.1$)	7.44	.1207	.332	6.00	.1058	.305	6.46	.525	.866	7.53	.536	.882
Ent-LS	3.16	.1116	.321	3.69	.1037	.308	3.17	.520	.862	8.58	.539	.885
CC-LS	1.48	.1111	.296	1.69	.1016	.279	4.15	.522	.873	2.54	.529	.880
LS-TS (ours)	1.49	.1111	.296	1.62	.1016	.279	4.16	.522	.873	2.65	.529	.881

corresponds to evaluating a proper scoring rule against a *degenerate* distribution δ_{y^*} that ignores genuine uncertainty; our methods replace this with the true ambiguity-aware distribution. This connection also clarifies why the *same* architectural class (temperature scaling) can be correct or incorrect depending solely on the target: the parameterization is not the bottleneck.

8.0.0.4 Extension to regression and ordinal settings.: While this paper focuses on classification, the core argument—that post-hoc calibration should target the correct marginal distribution over labels rather than a single voted label—applies equally to regression tasks where annotation ambiguity takes the form of disagreement over continuous or

ordinal outcomes. In such settings, the annotator distribution $\pi(\cdot | x)$ can be estimated from repeated annotations and used to define soft regression targets, replacing the typical mean-squared-error loss with a proper scoring rule against the empirical outcome distribution. We leave the development of ambiguity-aware post-hoc calibration for regression as future work.

8.0.0.5 Limitations.: Full ambiguity-aware calibration (including Dirichlet-Soft) requires multi-annotator labels at calibration time. When only voted labels are available, LS-TS remains useful but does not fully match Dirichlet-Soft, particularly on datasets where model confidence diverges from annotation ambiguity (e.g., ISIC 2019 ViT-S/16).

TABLE 6: **Multi-seed stability — all benchmarks and architectures.** Mean $\pm$ std of ECE_{true} (%) over 5 seeds. For CIFAR-10H and ChaosNLI the seed varies the cal/test split; for ISIC 2019 and DermaMNIST it varies the synthetic annotation sample. Standard deviations are small relative to the ECE gaps in all cases.

Dataset	Arch	TS	SLTS	Dir-Soft	IR-Soft
CIFAR-10H	ResNet-50	4.25 $\pm$ 0.25	1.61 $\pm$ 0.12	1.37 $\pm$ 0.09	0.89 $\pm$ 0.12
CIFAR-10H	ViT-B/16	4.39 $\pm$ 0.20	0.83 $\pm$ 0.14	0.74 $\pm$ 0.07	0.88 $\pm$ 0.16
ChaosNLI	RoBERTa-L	11.54 $\pm$ 0.54	3.85 $\pm$ 0.34	2.90 $\pm$ 0.14	2.40 $\pm$ 0.18
ChaosNLI	DeBERTa-v3	12.73 $\pm$ 0.60	4.41 $\pm$ 0.54	3.80 $\pm$ 0.35	2.61 $\pm$ 0.41
ISIC 2019	ENet-B4	18.84 $\pm$ 0.17	9.80 $\pm$ 0.22	8.45 $\pm$ 0.20	1.79 $\pm$ 0.16
ISIC 2019	ViT-S/16	16.93 $\pm$ 0.16	7.30 $\pm$ 0.13	6.20 $\pm$ 0.19	2.05 $\pm$ 0.13
DermaMNIST	ResNet-18	22.43 $\pm$ 0.28	5.08 $\pm$ 0.25	3.53 $\pm$ 0.39	2.33 $\pm$ 0.17
DermaMNIST	ViT-S/16	24.81 $\pm$ 0.08	3.61 $\pm$ 0.28	3.11 $\pm$ 0.21	2.09 $\pm$ 0.14




Fig. 9: **Robustness of ambiguity-aware calibration.** Rankings are stable across random seeds, calibration-set sizes, and annotation counts. Full results for all architectures and benchmarks are in Appendices .014, .012, and .013.

ISIC 2019 and DermaMNIST rely on synthetic annotators derived from clinician-reported agreement. Crucially, this is a *class-conditional* model: annotation confusion depends only on the consensus class label, not on instance-level visual difficulty. Real multi-reader medical datasets with per-image reader distributions would provide stronger validation and would capture the instance-level variation our synthetic model does not; concrete candidates include VinDr-CXR [40] (3 independent radiologists per chest X-ray, 15k images) and CheXpert [41] (5 independent radiologist annotations on the test set), both of which provide true per-image annotator distributions directly amenable to our framework. Conclusions about absolute ECE magnitudes on these two datasets should be interpreted with this in mind; the relative ranking of methods is expected to be robust to the specific confusion matrix used, as confirmed by the multi-seed annotation-generation ablation (Appendix .014).

9 CONCLUSION

We formalized the problem of confidence calibration under ambiguous ground truth and showed that standard post-hoc calibrators fitted on voted labels are systematically miscalibrated against the underlying ambiguous label distribution. Our flagship ambiguity-aware method, **Dirichlet-Soft**, requires no retraining and achieves the best or near-best Brier score and NLL across most benchmarks and architectures (best NLL in 7 of 8 settings; best or tied Brier in 7 of 8 settings). On CIFAR-10H and ChaosNLI (benchmarks with real per-instance multi-annotator distributions), it reduces true-label ECE by 59–86% relative to Temperature Scaling; ISIC 2019 and DermaMNIST results (with synthetic annotators) provide supportive evidence across medical

imaging tasks. The key insight is that the calibration target, not method capacity, is the limiting factor, evidenced by three complementary ablations: Dirichlet calibration with voted-label targets (Dirichlet-Hard) consistently *underperforms* simple TS; per-instance temperature adaptation (ATS) trained with voted-label supervision yields no improvement over TS and sometimes degrades it; and the same global-temperature architecture as TS, when supplied with smoothed pseudo-targets (LS-TS), reduces ECE by 7–78%. Together these confirm that the calibration target, not architectural flexibility, is the limiting factor. We recommend **Dirichlet-Soft** as the default wherever multi-annotator data is available, and **LS-TS** as a practical alternative when only voted labels are present.

REFERENCES

- [1] B. Kompa, J. Snoek, and A. L. Beam, “Second opinion needed: Communicating uncertainty in medical machine learning,” *NPJ Digital Medicine*, vol. 4, no. 1, p. 4, 2021.
- [2] D. Feng, A. Harakeh, S. L. Waslander, and K. Dietmayer, “A review and comparative study on probabilistic object detection in autonomous driving,” *IEEE Transactions on Intelligent Transportation Systems*, vol. 23, no. 8, pp. 9961–9980, 2022.
- [3] Y. Ovadia, E. Fertig, J. Ren, Z. Nado, D. Sculley, S. Nowozin, J. V. Dillon, B. Lakshminarayanan, and J. Snoek, “Can you trust your model’s uncertainty? Evaluating predictive uncertainty under dataset shift,” in *Advances in Neural Information Processing Systems*, vol. 32, 2019.
- [4] M. Abdar, F. Pourpanah, S. Hussain, D. Rezazadegan, L. Liu, M. Ghavamzadeh, P. Fieguth, X. Cao, A. Khosravi, U. R. Acharya *et al.*, “A review of uncertainty quantification in deep learning: Techniques, applications and challenges,” *Information Fusion*, vol. 76, pp. 243–297, 2021.
- [5] D. Stutz, A. G. Roy, T. Matejovicova, P. Strachan, A. T. Cemgil, and A. Doucet, “Conformal prediction under ambiguous ground truth,” *Transactions on Machine Learning Research (TMLR)*, 2023, arXiv:2307.09302.

- [6] C. Guo, G. Pleiss, Y. Sun, and K. Q. Weinberger, "On calibration of modern neural networks," in *Proceedings of the 34th International Conference on Machine Learning (ICML)*. PMLR, 2017, pp. 1321–1330.
- [7] B. Zadrozny and C. Elkan, "Transforming classifier scores into accurate multiclass probability estimates," in *Proceedings of the 8th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 2002, pp. 694–699.
- [8] J. C. Platt, "Probabilistic outputs for support vector machines and comparisons to regularized likelihood methods," in *Advances in Large Margin Classifiers*. MIT Press, 1999, pp. 61–74.
- [9] M. Kull, M. P. Nieto, M. Kängsepp, T. Silva Filho, H. Song, and P. Flach, "Beyond temperature scaling: Obtaining well-calibrated multi-class probabilities with Dirichlet calibration," in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 32, 2019.
- [10] M. Minderer, J. Djolonga, R. Romijnders, F. A. Hubis, X. Zhai, N. Houlsby, D. Tran, and M. Lucic, "Revisiting the calibration of modern neural networks," in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 34, 2021.
- [11] Y. Bai, S. Mei, H. Wang, and C. Xiong, "Don't just blame overparametrization for over-confidence: Theoretical analysis of calibration in binary classification," in *Proceedings of the 38th International Conference on Machine Learning (ICML)*. PMLR, 2021, pp. 566–576.
- [12] C. Szegedy, V. Vanhoucke, S. Ioffe, J. Shlens, and Z. Wojna, "Rethinking the inception architecture for computer vision," in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2016, pp. 2818–2826.
- [13] R. Müller, S. Kornblith, and G. Hinton, "When does label smoothing help?" in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 32, 2019.
- [14] J. C. Peterson, R. M. Battleday, T. L. Griffiths, and O. Russakovsky, "Human uncertainty makes classification more robust," in *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, 2019, pp. 9617–9626.
- [15] Y. Nie, X. Zhou, and M. Bansal, "What can we learn from collective human opinions on natural language inference data?" in *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2020, pp. 9131–9143.
- [16] A. N. Uma, T. Fornaciari, D. Hovy, E. Pavlick, M. Poesio, and B. Plank, "Learning from disagreement: A survey," *Journal of Artificial Intelligence Research*, vol. 72, pp. 1385–1470, 2021.
- [17] F. Rodrigues and F. Pereira, "Deep learning from crowds," in *Proceedings of the 32nd AAAI Conference on Artificial Intelligence*, 2018.
- [18] J. Baan, W. Aziz, B. Plank, and R. Fernández, "Stop measuring calibration when humans disagree," in *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2022, pp. 1892–1915.
- [19] M. L. Gordon, M. S. Lam, J. S. Park, K. Patel, J. Hancock, T. Hashimoto, and M. S. Bernstein, "Jury learning: Integrating dissenting voices into machine learning models," in *Proceedings of the CHI Conference on Human Factors in Computing Systems*, 2022.
- [20] B. Plank, "The problem with human label variation: On ground truth in data, modeling and evaluation," *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2022.
- [21] U. Khurana, E. Nalisnick, A. Fokkens, and S. Swayamdipta, "Crowd-calibrator: Can annotator disagreement inform calibration in subjective tasks?" in *Proceedings of the Conference on Language Modeling (COLM)*, 2024.
- [22] S. Park, O. Bastani, J. Weimer, and I. Lee, "Calibrated prediction with covariate shift via unsupervised domain adaptation," in *International Conference on Artificial Intelligence and Statistics (AISTATS)*. PMLR, 2020, pp. 3219–3229.
- [23] X. Wang, M. Long, J. Wang, and M. Jordan, "Transferable calibration with lower bias and variance in domain adaptation," in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 33, 2020.
- [24] S. Thulasidasan, G. Chennupati, J. A. Bilmes, T. Bhattacharya, and S. Michalak, "On mixup training: Improved calibration and predictive uncertainty for deep neural networks," in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 32, 2019.
- [25] G. Hinton, O. Vinyals, and J. Dean, "Distilling the knowledge in a neural network," *arXiv preprint arXiv:1503.02531*, 2015.
- [26] T. Joy, F. Pinto, S.-N. Lim, P. H. S. Torr, and P. K. Dokania, "Sample-dependent adaptive temperature scaling for improved calibration," in *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 37, no. 12, 2023, pp. 14 919–14 926.
- [27] J. Nixon, M. W. Dusenberry, L. Zhang, G. Jerfel, and D. Tran, "Measuring calibration in deep learning," in *IEEE Conference on Computer Vision and Pattern Recognition Workshops (CVPRW)*, 2019.
- [28] K. He, X. Zhang, S. Ren, and J. Sun, "Deep residual learning for image recognition," in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2016, pp. 770–778.
- [29] A. Dosovitskiy, L. Beyer, A. Kolesnikov, D. Weissenborn, X. Zhai, T. Unterthiner, M. Dehghani, M. Minderer, G. Heigold, S. Gelly, J. Uszkoreit, and N. Houlsby, "An image is worth 16x16 words: Transformers for image recognition at scale," *International Conference on Learning Representations (ICLR)*, 2021.
- [30] Y. Liu, M. Ott, N. Goyal, J. Du, M. Joshi, D. Chen, O. Levy, M. Lewis, L. Zettlemoyer, and V. Stoyanov, "RoBERTa: A robustly optimized BERT pretraining approach," *arXiv preprint arXiv:1907.11692*, 2019.
- [31] P. He, J. Gao, and W. Chen, "DeBERTaV3: Improving DeBERTa using ELECTRA-style pre-training with gradient-disentangled embedding sharing," in *International Conference on Learning Representations (ICLR)*, 2023.
- [32] P. Tschandl, C. Rosendahl, and H. Kittler, "The HAM10000 dataset, a large collection of multi-source dermatoscopic images of common pigmented skin lesions," *Scientific Data*, vol. 5, p. 180161, 2018.
- [33] M. Combalia, N. C. F. Codella, V. Rotemberg, B. Helba, V. Vilaplana, O. Reiter, C. Carrera, A. Barreiro, A. C. Halpern, S. Puig, and J. Malvehy, "BCN20000: Dermoscopic lesions in the wild," *arXiv preprint arXiv:1908.02288*, 2019.
- [34] A. P. Dawid and A. M. Skene, "Maximum likelihood estimation of observer error-rates using the EM algorithm," *Applied Statistics*, vol. 28, no. 1, pp. 20–28, 1979.
- [35] Y. Liu, A. Jain, C. Eng, D. H. Way, K. Lee, P. Bui, K. Kanada, G. de Oliveira Marinho, J. Gallegos, S. Gabriele, V. Gupta, N. Singh, V. Natarajan, R. Hofmann-Wellenhof, G. S. Corrado, L. H. Peng, D. R. Webster, D. Ai, S. Huang, Y. Liu, R. C. Dunn, and D. Coz, "A deep learning system for differential diagnosis of skin diseases," *Nature Medicine*, vol. 26, pp. 900–908, 2020.
- [36] M. Tan and Q. V. Le, "EfficientNet: Rethinking model scaling for convolutional neural networks," in *Proceedings of the 36th International Conference on Machine Learning (ICML)*, 2019, pp. 6105–6114.
- [37] H. A. Haenssle, C. Fink, R. Schneiderbauer, F. Toberer, T. Buhl, A. Blum, A. Kalloo, A. B. H. Hassen, L. Thomas, A. Enk, W. Stolz, and Reader Study Level-I and Level-II Investigators, "Man against machine: Diagnostic performance of a deep learning convolutional neural network for dermoscopic melanoma recognition in comparison to 58 dermatologists," *Annals of Oncology*, vol. 29, no. 8, pp. 1836–1842, 2018.
- [38] J. Yang, R. Shi, D. Wei, Z. Liu, L. Zhao, B. Ke, H. Pfister, and B. Ni, "MedMNIST v2: A large-scale lightweight benchmark for 2D and 3D biomedical image classification," *Scientific Data*, vol. 10, no. 1, p. 41, 2023.
- [39] T. Gneiting and A. E. Raftery, "Strictly proper scoring rules, prediction, and estimation," *Journal of the American Statistical Association*, vol. 102, no. 477, pp. 359–378, 2007.
- [40] H. Q. Nguyen, K. Lam, L. T. Le, H. H. Pham, D. Q. Tran, D. B. Nguyen, D. D. Le, C. M. Pham, H. T. T. Tong, D. H. Dinh et al., "VinDr-CXR: An open dataset of chest X-rays with radiologist's annotations," *Scientific Data*, vol. 9, no. 1, p. 429, 2022.
- [41] J. Irvin, P. Rajpurkar, M. Ko, Y. Yu, S. Cioffi, C. Chute, B. Kim, M. P. Lungren, A. Y. Ng, and P. Rajpurkar, "CheXpert: A large chest radiograph dataset with uncertainty labels and expert comparison," in *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 33, 2019, pp. 590–597.

APPENDIX

All four benchmarks in this paper are appropriate for studying calibration under ambiguous ground truth, but they exhibit ambiguity in slightly different ways.

.0.0.1 CIFAR-10H: Peterson et al. [14] introduced CIFAR-10H by collecting a full distribution of human labels for each image in the CIFAR-10 test set, explicitly to capture *human perceptual uncertainty*. The ambiguous examples are typically low-resolution images whose visual evidence supports multiple plausible classes for human observers. In this benchmark, the label distribution is therefore directly

observed from repeated human annotation, making it a clean testbed for ambiguity-aware calibration.

0.0.2 ChaosNLI: Nie et al. [15] introduced ChaosNLI to study *collective human opinions* on natural language inference. Each example receives 100 human labels, and the paper reports that substantial disagreement is common rather than exceptional. This ambiguity is semantic rather than perceptual: differences in pragmatic assumptions, underspecification, and sentence interpretation lead reasonable annotators to assign different NLI labels to the same premise-hypothesis pair.

0.0.3 ISIC 2019: ISIC 2019 is a multiclass dermatoscopic lesion diagnosis benchmark assembled from several clinical sources, including HAM10000 [32] and BCN20000 [33]. The task itself is ambiguity-prone because several lesion categories share overlapping visual morphology, and the challenge was designed to reflect the more realistic task of differential diagnosis rather than a simpler benign-versus-malignant decision. Because the public challenge release does not provide per-image reader distributions, we model ambiguity using synthetic annotators calibrated to dermatologist confusion patterns and reader agreement reported in Liu et al. [35]. This preserves the core property relevant to our study: in dermatology, a single image can support multiple clinically plausible labels.

0.0.4 DermaMNIST: DermaMNIST [38] is a standardized 28×28 version of the dermatology benchmark derived from HAM10000 [32]. HAM10000 consists of dermatoscopic images of common pigmented lesions collected in clinical practice, where diagnosis already involves fine-grained distinctions among visually similar categories such as melanoma, nevi, and benign keratoses. DermaMNIST inherits this medical ambiguity from HAM10000 and, because it is aggressively downsampled for lightweight benchmarking, removes additional visual detail that can further increase label uncertainty. As with ISIC 2019, we therefore evaluate calibration against an ambiguity-aware label distribution rather than treating the majority diagnosis as uniquely correct.

Let $\mathcal{D}_{\text{amb}} = \{(x_i, y_i^*) : y_i^* = c, x_i \in \mathcal{C}\}$ be the subset of calibration examples from the ambiguous cluster $\mathcal{C}$, where the voted label is always class c . Denote $p_i = \text{softmax}(z_i)_c$.

The TS loss on $\mathcal{D}_{\text{amb}}$ is $\mathcal{L}_{\text{TS}}(T; \mathcal{D}_{\text{amb}}) = -|\mathcal{D}_{\text{amb}}|^{-1} \sum_i \log \text{softmax}(z_i/T)_c$. As $T \rightarrow 0^+$, $\text{softmax}(z_i/T)_c \rightarrow 1$ (assuming $z_{ic} > z_{ik}$ for all $k \neq c$), so $\mathcal{L}_{\text{TS}} \rightarrow 0$. Hence $\partial \mathcal{L}_{\text{TS}} / \partial T > 0$ for all $T > 0$: decreasing T always decreases the voted-label loss on the ambiguous cluster. The globally optimal T_{TS}^* balances this downward pull from ambiguous examples against the upward pull from unambiguous examples, but $T_{\text{TS}}^* \leq T_{\text{no-amb}}^*$ (the optimal temperature without the ambiguous cluster): including ambiguous examples drives the optimum downward.

The soft-label loss on $\mathcal{D}_{\text{amb}}$ is $\mathcal{L}_{\text{soft}}(T; \mathcal{D}_{\text{amb}}) = -|\mathcal{D}_{\text{amb}}|^{-1} \sum_i \sum_k \hat{\pi}_k(x_i) \log \text{softmax}(z_i/T)_k$, with $\hat{\pi}_c(x_i) = q < 1$. The minimum-loss T_{soft}^* satisfies $\text{softmax}(z_i/T_{\text{soft}}^*)_c \approx q$, requiring $T_{\text{soft}}^* > T_{\text{TS}}^*$ whenever $p_i > q$. Hence $T_{\text{TS}}^* < T_{\text{soft}}^*$ regardless of whether either optimum exceeds 1. In the special case where the model is already overconfident on the ambiguous cluster ($p_i \gg q$), this gap is large enough that $T_{\text{TS}}^* < 1$; in practice both optima exceed 1 (since the model is also underconfident on unambiguous examples),

but the inequality $T_{\text{TS}}^* < T_{\text{soft}}^*$ is observed consistently across all experimental settings. $\square$

We formalize the argument given in the proof sketch in the main text.

0.0.5 Setup: Let $p = \hat{p}_{\hat{c}}(x)$ denote the model's predicted confidence for the top class $\hat{c}$ and let $q = \pi_{\hat{c}}(x) \in (0, 1]$ denote the true annotator probability for that class. We assume the model is trained against voted labels, so p is approximately constant across examples with the same voted label regardless of their annotation entropy $H(x)$.

0.0.6 Voted-label calibration error contribution: For example x_i in confidence bin B_b (where $\bar{p}(B_b) \approx p$), the per-example voted-label calibration error is $|p - \mathbf{1}[\hat{c} = y^*]|$. Since $\hat{c} = y^*$ for correctly-classified ambiguous examples (the model predicts the majority class), this contribution is approximately $|p - 1|$ for examples where the model makes the right hard prediction, independent of $H(x)$.

0.0.7 True-label calibration error contribution: The per-example true-label calibration error is $|p - \pi_{\hat{c}}(x)| = |p - q|$.

0.0.8 Excess true-label error relative to voted-label evaluation: The per-example excess contribution is:

$$\delta(x) = |p - q| - |p - \mathbf{1}[\hat{c} = y^*]|. \quad (6)$$

When the model is correctly calibrated to voted labels ($p \approx \mathbf{1}[\hat{c} = y^*]$ in expectation), the voted contribution is near zero: $|p - 1| \approx 0$ for high-accuracy bins and $|p - 0| = p$ for error bins. For unambiguous examples ($H(x) = 0$), $q = \pi_{y^*} = 1$, so $\delta(x) = |p - 1| - |p - 1| = 0$.

0.0.9 Monotonicity: As $H(x)$ increases, $q = \pi_{y^*}(x)$ decreases below 1 (more probability mass shifts to non-majority classes). Specifically, for a K -class problem with uniform annotator distribution at maximum entropy, $q = 1/K$. The true-label contribution $|p - q|$ increases strictly as q decreases from 1 (since $p > q$ for a well-trained model predicting the majority class with confidence exceeding the majority probability). The voted-label contribution $|p - \mathbf{1}[\hat{c} = y^*]|$ is unaffected by $H(x)$. Therefore $\delta(x) = |p - q| - |p - \mathbf{1}[\hat{c} = y^*]|$ is non-decreasing in $H(x)$, and strictly increasing whenever $p > q > 0$.

0.0.10 Aggregate monotonicity: The aggregate excess true-label error relative to voted-label evaluation equals $\sum_b \frac{|B_b|}{n} \sum_{i \in B_b} \delta(x_i) / |B_b|$. Since $\delta(x_i)$ is non-decreasing in $H(x_i)$ and bins contain a mixture of examples, this discrepancy grows with the mean annotation entropy of the test set. $\square$

Remark 1 (Independence assumption). *Assumption (i) (that $\hat{p}_{\hat{c}}(x)$ is approximately independent of $H(x)$) may be violated in practice if harder examples (high $H(x)$) also have lower model confidence. When the two are positively correlated, the per-example discrepancy term $\delta(x)$ grows even faster with $H(x)$ than the proof suggests, so Proposition 2 remains valid and the monotonicity is if anything conservative.*

0.0.11 Method: When individual annotation records are available rather than pre-aggregated label distributions, MCTS draws S samples $a_{is} \sim \hat{\pi}(x_i)$ per example and minimizes

$$\mathcal{L}_{\text{MCTS}}(T) = -\frac{1}{nS} \sum_{i=1}^n \sum_{s=1}^S \log \text{softmax}(z_i/T)_{a_{is}}.$$




Fig. 10: **Calibration-set-size ablation — CIFAR-10H.** Left: ResNet-50. Right: ViT-B/16. TS’s ECE is flat across all calibration set sizes; ambiguity-aware methods converge within 5–10%.




Fig. 11: **Calibration-set-size ablation — ChaosNLI.** Left: RoBERTa-Large. Right: DeBERTa-v3-base. NLI models show the same pattern: TS is flat near 10–13% while soft methods stabilise near 3–4%.

MCTS is a strictly more general formulation than SLTS: it applies even when only raw annotation samples are stored rather than pre-computed frequency vectors.

.0.0.12 A single annotation suffices.: The key practical implication of MCTS is that *even* $S = 1$ (a single randomly drawn annotation per calibration example) already achieves ECE_{true} of 1.53%, matching SLTS to within rounding error and reducing TS’s 4.29% by 64%. This means that any dataset where each example has been annotated even once by a randomly selected annotator (rather than deterministically assigned the majority label) is immediately amenable to strong ambiguity-aware calibration. Formally, since $\mathbb{E}_{\hat{y} \sim \hat{\pi}}[\text{CE}(\text{softmax}(z/T), \hat{y})] = \text{KL}(\hat{\pi} \parallel \text{softmax}(z/T)) + H(\hat{\pi})$, the MCTS objective converges to SLTS in expectation as $S \rightarrow \infty$, with variance $\propto 1/S$. Table 1 (main text) confirms the fast convergence: the ECE gap between $S = 1$ and $S \rightarrow \infty$ is only 0.02 pp; in practice $S = 20$ –50 is sufficient to reduce variance to negligible levels.

We fix the annotation count at the full observed value and vary the fraction of calibration examples used, from 5% to 100% (stratified subsampling by class). Results are shown for all four benchmarks and eight architectures.

.0.0.13 Key findings.: Across all four benchmarks and eight architectures, TS’s ECE does not decrease with more calibration data, confirming that the degradation is a target bias rather than a variance artifact. In contrast, SLTS achieves stable low ECE from as few as 5–10% of the calibration set, and IR-Soft and Dirichlet-Soft converge even faster on most benchmarks. This suggests that the annotation distribution (not the number of calibration examples) is the limiting factor for ambiguity-aware methods.

We study how performance depends on the number of annotations m per calibration image. We fix the calibration set size and the annotation generation seed, and sweep m from 1 to 50, resampling the soft targets for each value of m . Results are shown for CIFAR-10H, which provides up to ≈ 51




Fig. 12: **Calibration-set-size ablation — ISIC 2019.** Left: EfficientNet-B4. Right: ViT-S/16. IR-Soft achieves the lowest ECE across all calibration set sizes on both architectures.




Fig. 13: **Calibration-set-size ablation — DermaMNIST.** Left: ResNet-18. Right: ViT-S/16. Ambiguity-aware methods converge quickly; TS remains flat regardless of calibration set size.

real annotations per image; we subsample from the full pool.

.0.0.14 Key findings.: Ambiguity-aware calibration is robust to annotation count once $m \geq 2$: SLTS, Dirichlet-Soft, and IR-Soft achieve near-minimal ECE with only a handful of annotations per image and do not improve substantially beyond $m = 5$. LS-TS converges quickly (by $m \approx 4$) because the global smoothing parameter ε stabilises once each image has at least a few annotations. Standard Temperature Scaling shows the opposite trend: its ECE_{true} *increases* with m because a higher temperature is required to match the more diffuse annotation distribution, exposing the fundamental target mismatch.

For CIFAR-10H and ChaosNLI we repeat the evaluation with five random seeds (42–46), each producing a different stratified 50/50 cal/test split. For ISIC 2019 and DermaMNIST the val/test split is predetermined, so we instead vary the annotation generation seed (42–46), testing sensitivity to the particular synthetic annotation sample. Table 6 (main text, Robustness section) reports mean $\pm$ std of ECE_{true} across seeds for all eight (dataset, arch) combinations.

Across all benchmarks and architectures, standard deviations are small relative to the ECE gaps: TS vs. ambiguity-aware methods differ by 2–10 $\times$ the standard deviation, confirming statistical robustness of the results in Tables 2–4. Bar charts for all architectures are shown below.

Four-method summaries for CIFAR-10H (ResNet-50) and ChaosNLI (RoBERTa-Large) are shown in Figures 4 and 5 in the main text. Here we provide the four-method summary panels and full panels for all remaining architectures and datasets. All panels share the same layout: top row = voted-label baselines (Uncal, TS, Platt, Dir.-Hard); middle row = annotation-free (LS-TS) and temperature-based soft methods (SLTS, MCTS, SoftPlatt, VS); bottom row = non-parametric/matrix soft methods (Dir.-Soft, IR-Soft). Red shading = overconfident; green = underconfident.




Fig. 14: Annotation-count ablation — CIFAR-10H (both architectures). ECE_{true} vs. annotations per image m for ResNet-50 (left) and ViT-B/16 (right). Ambiguity-aware methods are stable for $m \geq 2$; TS rises with m due to target mismatch; LS-TS converges by $m \approx 4$.




Fig. 15: Multi-seed bar charts — CIFAR-10H. Left: ResNet-50. Right: ViT-B/16.

- 1 CIFAR-10H
- 2 ChaosNLI
- 3 ISIC 2019
- 4 DermaMNIST

The diagrams confirm the quantitative results in Tables 2–4: TS reduces overconfidence relative to *Uncal* but leaves a large residual gap because it targets voted labels. Dirichlet-Hard performs similarly to TS and sometimes worse, confirming that the problem lies in the target, not model capacity. All soft methods substantially close the gap, with IR-Soft and Dir-Soft achieving the most uniform residuals across all benchmarks.

5 Construction

Because the full per-image reader annotations of Liu et al. [35] are not publicly available, we construct a synthetic annotator model using a clinically calibrated 8×8 confusion matrix C , where entry C_{ij} gives the probability that a board-certified dermatologist labels an image as class j when the consensus (majority-vote) diagnosis is class i . This confusion-matrix annotator model follows the framework of Dawid and Skene [34], in which each annotator’s behaviour is characterised by a class-conditional label-flipping distribution; it is widely used in crowdsourcing and multi-annotator learning [16], [17].

5.0.1 Calibration sources.: Diagonal entries (per-condition agreement rates) are set to match the per-condition inter-reader agreement reported in Liu et al. [35] (Table 2, nine board-certified dermatologists reading 963 clinical images) and corroborated by Haenssle et al. [37]. Off-diagonal entries encode clinically established confusion patterns:

- **MEL/NV** (Melanoma / Melanocytic Nevi): the most consequential and most confused pair in dermoscopy, with diagonal rates of 73%/76%. Off-diagonal entries




Fig. 16: Multi-seed bar charts — ChaosNLI. Left: RoBERTa-Large. Right: DeBERTa-v3-base.




Fig. 17: Multi-seed bar charts — ISIC 2019. Left: EfficientNet-B4. Right: ViT-S/16. Seed varies the annotation generation; val/test split is fixed.

$C_{\text{MEL,NV}} = 0.14$ and $C_{\text{NV,MEL}} = 0.15$ reflect the well-documented difficulty of distinguishing early melanoma from benign nevi [37].

- **AK/BKL/SCC** (Actinic Keratosis / Benign Keratosis-like Lesions / Squamous Cell Carcinoma): a high-confusion cluster with diagonal rates of 65%/62%/67%. The $\text{AK} \leftrightarrow \text{SCC}$ confusion ($C_{\text{AK,SCC}} = 0.16$, $C_{\text{SCC,AK}} = 0.18$) is clinically significant because untreated AK can progress to SCC. BKL shares morphological features with both conditions.
- **DF/VL** (Dermatofibroma / Vascular Lesion): the easiest to distinguish, with diagonal rates of 87%/91%.

The mean diagonal agreement is $\bar{C}_{ii} = 75\%$, matching the aggregate inter-reader agreement reported by Liu et al. [35] exactly.

5.0.2 Simulation procedure.: For each test image x_i with consensus label y_i , we independently draw $m = 9$ synthetic annotations from the categorical distribution defined by row y_i of C :

$$a_i^{(k)} \sim \text{Categorical}(C_{y_i, \cdot}), \quad k = 1, \dots, m,$$

and set the empirical label distribution to $\pi(x_i) = \frac{1}{m} \sum_{k=1}^m \mathbf{e}_{a_i^{(k)}}$. This directly mirrors the $m = 9$ board-certified dermatologist annotation protocol of Liu et al. [35], whose reader study we cannot fully replicate due to data unavailability. The resulting mean annotation entropy over the ISIC 2019 test set is $\bar{H} = 0.66$, reflecting the genuine diagnostic difficulty of the 8-class task.

Remark 2 (Class-conditional limitation). *This synthetic annotator model is deliberately class-conditional: every image with the same consensus class y_i is assigned the same row $C_{y_i, \cdot}$ of the confusion matrix, regardless of its individual visual content. In practice, images within the same consensus class will differ in visual ambiguity, leading to instance-level variation in annotator disagreement that this model does not capture. Our experiments therefore evaluate calibration under a controlled form of clinically-informed ambiguity rather than true per-instance annotator distributions.*




Fig. 18: Multi-seed bar charts — DermaMNIST. Left: ResNet-18. Right: ViT-S/16. Seed varies the annotation generation; val/test split is fixed.




Fig. 19: Reliability diagrams (summary) — CIFAR-10H ViT-B/16.

.6 Visualisation

Figure 33 shows the confusion matrix as a heatmap. The MEL/NV cluster (orange dashed box) and the AK/BKL/SCC high-confusion cluster (purple dotted highlights) are visually prominent. DF and VL have near-perfect diagonal entries and negligible off-diagonal mass.

.7 Full confusion matrix

Table 7 lists the exact numerical entries of the confusion matrix.

Proposition 2 states that, under simplifying assumptions, the pointwise true-label calibration error of a voted-label-calibrated model is non-decreasing in the annotation entropy $H(x)$. Figure 3 provides empirical support: we bin test examples by their normalized annotation entropy $H(x)/\log K$ and plot the mean absolute error $|\hat{p}_\varepsilon(x) - \pi_\varepsilon(x)|$ for TS and SLTS.

For CIFAR-10H (both architectures), TS error rises from $\approx 2\%$ in the near-unambiguous bin ($H/\log K \approx 0.02$) to $\approx 18\%$ in the high-ambiguity bin ($H/\log K \approx 0.48$), a clear monotone increase consistent with the proposition. SLTS also increases (the soft target itself becomes harder to match at high entropy) but does so more moderately. For ChaosNLI (DeBERTa-v3), which operates in a higher-entropy regime ($H/\log K \in [0.2, 0.9]$), TS again rises at the extremes while SLTS attenuates the error at moderate entropy levels. The crossover in the low-entropy bins reflects that SLTS over-smooths near-unambiguous examples (where voted-label and soft-label targets nearly coincide), consistent with the slight performance difference between TS and SLTS on unambiguous subsets.

See Figure 3 in the main text (Theory section) for the empirical validation plot.

We compare LS-TS against three simpler annotation-free smoothing baselines that share the same single-temperature family but differ in how the pseudo-target $\tilde{\pi}_i^{\text{LS}}$ is constructed:

- **Fixed-LS** ($\varepsilon=0.1$): fixed data-independent smoothing weight, the standard label-smoothing value used in training [12].




Fig. 20: Reliability diagrams (all methods) — CIFAR-10H ResNet-50 (ECE_{true}).




Fig. 21: Reliability diagrams (all methods) — CIFAR-10H ViT-B/16 (ECE_{true}).

- **Ent-LS**: per-instance $\varepsilon_i = H(\hat{p}(x_i))/\log K$, using normalised prediction entropy as a proxy for annotator disagreement.
- **CC-LS**: per-class $\varepsilon_k = \text{mean}_i: y_i^* = k(1 - \hat{p}_{y^*}(x_i))$, accommodating class-level variation in model confidence (same idea as Vector Scaling applied to the smoothing weight).

All methods are evaluated on CIFAR-10H and ChaosNLI, the two benchmarks with real multi-annotator distributions, using the same experimental protocol as the main paper (seed 42, 50/50 stratified split). Table 5 (main text, Methods section) reports the full comparison.

Fixed-LS with $\varepsilon=0.1$ consistently *underperforms* or even *degrades* below TS (e.g., ECE 7.44% vs. 3.90% on CIFAR-10H ResNet-50), confirming that an arbitrary fixed smoothing weight is insufficient; the data-driven estimate of $\bar{\varepsilon}$ is critical. Ent-LS (per-instance entropy) improves over TS on most settings but is inconsistent: it works well on ChaosNLI RoBERTa-L (3.17%) but collapses on DeBERTa-v3 (8.58%), suggesting that prediction entropy is a noisy proxy for annotator ambiguity when models are already strongly confident. CC-LS (per-class ε) achieves results very close to LS-TS on all four settings (within 0.1 pp ECE), confirming that the global mean is a good summary of per-class behaviour for these benchmarks; the slight advantage of CC-LS on DeBERTa-v3 (2.54% vs. 2.65%) is within the single-seed




Fig. 22: Reliability diagrams (summary) — ChaosNLI DeBERTa-v3-base.




Fig. 23: Reliability diagrams (all methods) — ChaosNLI RoBERTa-Large (ECE_{true}). NLI models are severely overconfident before calibration; ambiguity-aware methods substantially correct this.

variance reported in Appendix .014.

A natural question is whether making the calibration map *adaptive* (predicting a per-instance temperature $T(x_i)$ rather than a global T^*) can close the gap between voted-label calibration and ambiguity-aware calibration, without requiring annotator data. We implement **ATS** [26]: a linear layer maps four logit-derived features $\phi(z_i) = [\max_k z_{ik}, H(\text{softmax}(z_i))/\log K, p_{(1)} - p_{(2)}, p_{(1)}]$ to $\log T_i$ via $T_i = \text{softplus}(w^\top \phi(z_i) + b) + 0.1$, trained with voted-label NLL identical to TS. The four features capture prediction scale, uncertainty, margin, and confidence. We apply ℓ_2 regularisation ($\lambda = 10^{-3}$) to avoid overfitting on small calibration sets.

7.0.1 Finding.: ATS fails to improve over TS in all 8 settings and in several cases degrades it (ECE rises from 11.63% to 13.09% on DeBERTa-v3; from 16.59% to 17.53% on ISIC ViT-S/16). The per-instance temperature learned from logit features adapts to *model uncertainty*, not to *annotator disagreement*; these two quantities are loosely correlated at best (cf. Figure 3). In contrast, LS-TS, which uses the same global-temperature architecture as TS but simply shifts the calibration target toward a smoothed distribution, reduces ECE by 7–78% across all 8 settings without any annotator data. This confirms that the calibration target is the critical design axis: added architectural flexibility without the right supervision signal yields no benefit.

7.0.2 CIFAR-10H models.: *ResNet-50*: ImageNet-1k weights (ResNet50_Weights.IMAGENET1K_V1); FC layer replaced with 2048×10 linear; AdamW, $\text{lr} = 10^{-4}$, weight decay 10^{-4} , batch 128, cosine annealing, 30 epochs, 224×224 . *ViT-B/16*: ImageNet-21k weights; classification head replaced;




Fig. 24: Reliability diagrams (all methods) — ChaosNLI DeBERTa-v3-base (ECE_{true}).




Fig. 25: Reliability diagrams (summary) — ISIC 2019 EfficientNet-B4.

AdamW, $\text{lr} = 5 \times 10^{-5}$, weight decay 10^{-4} , batch 64, cosine annealing, 20 epochs, 224×224 .

7.0.3 ChaosNLI models.: *RoBERTa-Large*: roberta-large-mnli from HuggingFace, pre-fine-tuned on MultiNLI [30]; no additional fine-tuning. *DeBERTa-v3-base*: cross-encoder/nli-deberta-v3-base from HuggingFace [31]; no additional fine-tuning. Both models use the combined ChaosNLI SNLI+MNLI subset ($n = 3,113$), split 50/50 into calibration and test (stratified, seed 42).

7.0.4 ISIC 2019 models.: *EfficientNet-B4*: ImageNet-1k weights; stratified 70/15/15 split; class-weighted cross-entropy (NV: $\approx 67\%$ of training images); AdamW, $\text{lr} = 3 \times 10^{-4}$, weight decay 10^{-4} , 2-epoch linear warm-up + cosine decay, 20 epochs, 380×380 . *ViT-S/16*: ImageNet-21k weights; same split and loss weighting; AdamW, $\text{lr} = 10^{-4}$, weight decay 10^{-4} , 2-epoch warm-up + cosine decay, 20 epochs, 224×224 .

7.0.5 DermaMNIST models.: *ResNet-18*: ImageNet-1k weights, class-weighted cross-entropy, AdamW, $\text{lr} = 10^{-4}$, 30 epochs. *ViT-S/16*: ImageNet-21k weights, same training protocol, 224×224 .

7.0.6 Calibration optimizers.: TS, SLTS: LBFGS, $\text{lr} = 0.1$, 500 iterations, tolerance 10^{-9} (gradient), 10^{-11} (loss). Platt/VS: Adam, $\text{lr} = 0.01/0.05$, weight decay 10^{-4} , 2000 steps.

7.0.7 ECE bins.: $B = 15$ equal-width bins in $[0, 1]$; bins with fewer than 1 example are excluded.

7.0.8 Compute.: CIFAR-10H fine-tuning: ≈ 20 min on one NVIDIA A100. All calibration methods: < 60 s on CPU after logit caching.




Fig. 26: Reliability diagrams (summary) — ISIC 2019 ViT-S/16.




Fig. 27: Reliability diagrams (all methods) — ISIC 2019 EfficientNet-B4 (ECE_{true}).




Fig. 28: Reliability diagrams (all methods) — ISIC 2019 ViT-S/16 (ECE_{true}).




Fig. 29: Reliability diagrams (summary) — DermaMNIST ResNet-18.




Fig. 30: Reliability diagrams (summary) — DermaMNIST ViT-S/16.




Fig. 31: Reliability diagrams (all methods) — DermaMNIST ResNet-18 (ECE_{true}).




Fig. 32: Reliability diagrams (all methods) — DermaMNIST ViT-S/16 (ECE_{true}).

TABLE 7: Numerical values of the inter-reader confusion matrix for ISIC 2019. Row = consensus (majority-vote) label; column = annotator's label. Each row sums to 1. Diagonal entries (bold) are per-condition agreement rates, calibrated to Liu et al. [35] Table 2; mean diagonal agreement $\bar{C}_{ii} = 75\%$.

	MEL	NV	BCC	AK	BKL	DF	VL	SCC
MEL	0.73	0.14	0.02	0.03	0.08	0.00	0.00	0.00
NV	0.15	0.76	0.01	0.01	0.06	0.01	0.00	0.00
BCC	0.02	0.01	0.81	0.05	0.07	0.01	0.01	0.02
AK	0.03	0.01	0.04	0.65	0.11	0.00	0.00	0.16
BKL	0.12	0.05	0.03	0.10	0.62	0.00	0.00	0.08
DF	0.01	0.02	0.02	0.01	0.02	0.87	0.03	0.02
VL	0.00	0.01	0.02	0.01	0.01	0.02	0.91	0.02
SCC	0.01	0.01	0.03	0.18	0.09	0.00	0.01	0.67

TABLE 8: **ATS vs. TS and LS-TS across all 8 settings (annotation-free methods only) — full results.** ECE_{true} (%) and Brier. ATS predicts a per-instance temperature from logit features but is trained with the same voted-label NLL as TS. **Bold** = best annotation-free method per column.

Method	C10H R50		C10H ViT		NLI RoBERTa		NLI DeBERTa		ISIC ENet		ISIC ViT		Derm R18		Derm ViT	
	ECE	Br	ECE	Br	ECE	Br	ECE	Br	ECE	Br	ECE	Br	ECE	Br	ECE	Br
TS	3.90	.113	4.01	.104	10.55	.537	11.63	.549	18.54	.560	16.59	.617	23.05	.686	24.36	.683
ATS	4.10	.113	4.10	.105	11.36	.538	13.09	.552	18.83	.560	17.53	.631	22.82	.687	25.35	.686
LS-TS	1.49	.111	1.62	.102	4.16	.522	2.65	.529	9.34	.528	15.49	.619	5.05	.630	7.36	.615




Fig. 33: **Inter-reader confusion matrix for ISIC 2019.** Entry C_{ij} is the probability that a dermatologist labels an image as class j given the consensus label is i . Diagonal entries (bold) are the per-condition agreement rates, calibrated to match Liu et al. [35] Table 2. Orange dashed box: MEL/NV most-confused pair. Purple dotted highlights: AK/BKL/SCC high-confusion cluster. Mean diagonal agreement $\bar{C}_{ii} = 75\%$.